

VIRTUALIZATION

OpenVZ

Checkpoints

- Introduction to virtualization
- Openvz running
- Container creation

What is Virtualization?

- Method of dividing computer resources
- Multiple isolated environments
- Types
 - Emulation
 - Paravirtualization
 - OS level virtualization

Operating System level virtualization

- Multiple isolated execution environments within a single OS kernel
- Doesn't allow to run different kernels from different OSs at the same time
- Best possible performance
- Examples: OpenVZ, Linux-VServer

Virtual Environment / Container

- Look and feel of separate physical server
- VE has its own
 - Process tree with init
 - File system
 - Network interfaces with IP addresses, Firewall rules, Routing tables

Virtual Environment

- Multiple VEs co-exist within a single physical server.
- Different VEs can run different Linux distros
- But, all VEs operate under same kernel

OpenVZ kernel

- Modified Linux based kernel which comes with following functionalities
 - Virtualization and Isolation
 - Enables many VEs within a single kernel
 - Resource management
 - Manages CPU, RAM, Disk space per VE basis
 - Checkpointing
 - Freezing a VE, Saving its state to disk with ability to unfreeze that state later

OpenVZ vs VMware/virtualbox

- Virtualbox - Hardware assisted
- OpenVZ - OS level virtualization

Demo on OpenVZ

- Vzctl
 - create
 - start
 - set
 - restart
 - status
 - enter
 - stop
- Vzlist ..

thanks

Virtualization and Isolation

- Each VE has its own set of resources provided by OS kernel
 - Files
 - System libraries, applications, virtualized /proc & /sys, virtualized lock ,.etc
 - users and groups
 - process tree
 - network
 - Allows VE to have its own IP addresses

Virtualization and Isolation

- Devices
 - a VE can be granted to access real device like network interfaces, serial ports
- IPC objects
 - shared memory
 - semaphores
 - messages

Resource Management

- Finite set of resources within a single kernel are shared among multiple VEs
- OpenVZ resource management subsystem consists of three components
 - Two-level disk quota
 - Fair CPU scheduler
 - UBC